

Importing Liabraries

```
In [1]: import pandas as pd
import numpy as np
```

Loading Dataset

```
In [3]: dataset=pd.read_csv(r'D:\csv files\Brain Stroke.csv')
dataset
```

Out[3]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_gl
0	Male	67.0	0	1	Yes	Private	Urban	
1	Male	80.0	0	1	Yes	Private	Rural	
2	Female	49.0	0	0	Yes	Private	Urban	
3	Female	79.0	1	0	Yes	Self-employed	Rural	
4	Male	81.0	0	0	Yes	Private	Urban	
...
4976	Male	41.0	0	0	No	Private	Rural	
4977	Male	40.0	0	0	Yes	Private	Urban	
4978	Female	45.0	1	0	Yes	Govt_job	Rural	
4979	Male	40.0	0	0	Yes	Private	Rural	
4980	Female	80.0	1	0	Yes	Private	Urban	

4981 rows × 11 columns



Preprocessing

```
In [6]: from sklearn.preprocessing import LabelEncoder
l1=LabelEncoder()
```


```
In [7]: dataset['gender']=l1.fit_transform(dataset['gender'])
dataset['ever_married']=l1.fit_transform(dataset['ever_married'])
dataset['work_type']=l1.fit_transform(dataset['work_type'])
dataset['Residence_type']=l1.fit_transform(dataset['Residence_type'])
dataset['smoking_status']=l1.fit_transform(dataset['smoking_status'])
```

In [8]: dataset

Out[8]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_glucose_level
0	1	67.0	0	1	1	1	1	
1	1	80.0	0	1	1	1	0	
2	0	49.0	0	0	1	1	1	
3	0	79.0	1	0	1	2	0	
4	1	81.0	0	0	1	1	1	
...
4976	1	41.0	0	0	0	1	0	
4977	1	40.0	0	0	1	1	1	
4978	0	45.0	1	0	1	0	0	
4979	1	40.0	0	0	1	1	0	
4980	0	80.0	1	0	1	1	1	

4981 rows × 11 columns



In [9]: `x=dataset.iloc[:, :-1].values`
`y=dataset.iloc[:, -1].values`

In [10]: x

Out[10]: `array([[1. , 67. , 0. , ..., 228.69, 36.6 , 1.],`
`[1. , 80. , 0. , ..., 105.92, 32.5 , 2.],`
`[0. , 49. , 0. , ..., 171.23, 34.4 , 3.],`
`...,`
`[0. , 45. , 1. , ..., 95.02, 31.8 , 3.],`
`[1. , 40. , 0. , ..., 83.94, 30. , 3.],`
`[0. , 80. , 1. , ..., 83.75, 29.1 , 2.]])`

In [11]: y

Out[11]: `array([1, 1, 1, ..., 0, 0, 0], dtype=int64)`

In [13]: `from sklearn.preprocessing import StandardScaler`
`sc=StandardScaler()`
`scale=sc.fit_transform(x)`

In [14]: scale

```
Out[14]: array([[ 1.18390850e+00,  1.04058433e+00, -3.26185770e-01, ...,
                2.72341090e+00,  1.19323816e+00, -3.53933192e-01],
                [ 1.18390850e+00,  1.61427033e+00, -3.26185770e-01, ...,
                -5.22766599e-04,  5.89389611e-01,  5.78839946e-01],
                [-8.44659868e-01,  2.46249882e-01, -3.26185770e-01, ...,
                 1.44852918e+00,  8.69221866e-01,  1.51161308e+00],
                ...,
                [-8.44659868e-01,  6.97311148e-02,  3.06573766e+00, ...,
                -2.42364234e-01,  4.86293516e-01,  1.51161308e+00],
                [ 1.18390850e+00, -1.50917344e-01, -3.26185770e-01, ...,
                -4.88199415e-01,  2.21189274e-01,  1.51161308e+00],
                [-8.44659868e-01,  1.61427033e+00,  3.06573766e+00, ...,
                -4.92415000e-01,  8.86371531e-02,  5.78839946e-01]])
```

In [15]: dataset['stroke'].value_counts()

```
Out[15]: 0    4733
         1     248
         Name: stroke, dtype: int64
```

```
In [16]: from imblearn.over_sampling import RandomOverSampler
         ros=RandomOverSampler()
         X,Y=ros.fit_resample(scale,y)
```

```
In [17]: from collections import Counter
         print(Counter(Y))
```

```
Counter({1: 4733, 0: 4733})
```

Splitting data into train and test set

```
In [18]: from sklearn.model_selection import train_test_split
         x_train,x_test,y_train,y_test=train_test_split(X,Y,test_size=0.2,random_state=9)
```

Logistic Regression

```
In [19]: from sklearn.linear_model import LogisticRegression
         lg=LogisticRegression()
         lg.fit(x_train,y_train)
```

```
Out[19]: LogisticRegression()
```

```
In [20]: y_pred=lg.predict(x_test)
         y_pred
```

```
Out[20]: array([1, 1, 1, ..., 1, 1, 1], dtype=int64)
```

Accuracy of Logistic Regression

```
In [23]: from sklearn.metrics import accuracy_score  
print(accuracy_score(y_pred,y_test)*100)
```

75.55438225976768

KNN

```
In [25]: from sklearn.neighbors import KNeighborsClassifier  
knn=KNeighborsClassifier(n_neighbors=5,metric='minkowski',p=2)  
knn.fit(x_train,y_train)
```

Out[25]: KNeighborsClassifier()

```
In [26]: y1_pred=knn.predict(x_test)  
y1_pred
```

Out[26]: array([1, 1, 1, ..., 0, 1, 1], dtype=int64)

Accuracy of KNN

```
In [27]: from sklearn.metrics import accuracy_score  
print(accuracy_score(y1_pred,y_test)*100)
```

92.18585005279832

In []: